

HLD Chapter 1: Distributed Systems Core Prerequisites & Latency Numbers

Core Learning Objective: Master the foundational theorems, network protocols, and hardware latency figures that govern every distributed system at planetary scale.

1. Latency Numbers Every Systems Engineer Must Know

Peter Norvig and Jeff Dean's famous hardware latency comparison demonstrates why network calls and disk reads require caching and asynchronous architectures:

Operation	Real Time	Scaled Human Equivalent (1 CPU cycle $\approx$ 1 sec)
L1 CPU Cache reference	0.5 ns	1 second
L2 CPU Cache reference	7 ns	14 seconds
Main RAM memory access	100 ns	3.3 minutes
SSD random read (NVMe)	16 μ s	5.5 hours
Rotational Disk Seek (HDD)	4 ms	1.3 months
Same Datacenter Roundtrip (LAN)	0.5 ms	5.7 days
Cross-Continent Roundtrip (SF to London)	150 ms	4.7 years!

IMPORTANT


Key Takeaway: Reading from RAM is $\sim 160\times$ faster than an SSD, and $\sim 40,000\times$ faster than a mechanical disk! Sending a network packet across the ocean is equivalent to waiting almost 5 years. This is why CDNs and distributed caches exist.

2. CAP Theorem vs PACELC Theorem

The Classic CAP Theorem

In any asynchronous distributed network subject to partitions (P), you can guarantee at most **two** of the following three properties:

- **Consistency (C)**: Every read receives the most recent write or an error (Linearizability).
- **Availability (A)**: Every non-failing node returns a non-error response, though not guaranteed to be the newest write.
- **Partition Tolerance (P)**: The system continues to operate despite arbitrary network drops between nodes.



NOTE

Since network partitions are an unavoidable physical reality in distributed networks, "**CA**" systems do not exist in distributed computing. Your true choice is **CP vs AP**.

The PACELC Theorem (The Complete Reality)

CAP only describes system behavior *during a network partition*. The **PACELC Theorem** extends this to normal operations:


- **If there is a Partition (P)**: How does the system trade off **Availability (A)** vs **Consistency (C)**?

- **Else (E):** When the network is normal, how does the system trade off **Latency (L)** vs **Consistency (C)**?

Examples:

- **PA/EL (Cassandra, DynamoDB):** On Partition $\rightarrow$ Available. Else $\rightarrow$ Low Latency (Eventual Consistency).
- **PC/EC (MongoDB, Spanner):** On Partition $\rightarrow$ Consistent. Else $\rightarrow$ Strong Consistency (higher latency due to sync replication).

3. Network Transport Protocols: TCP, UDP, HTTP/2, HTTP/3, gRPC & WebSockets



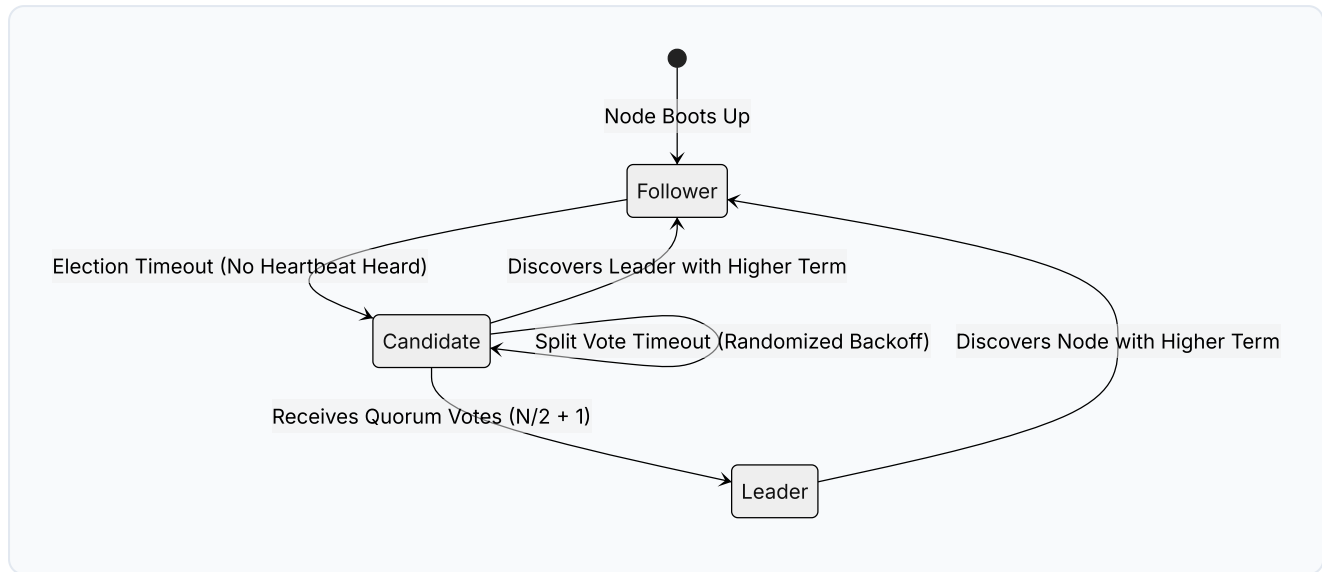
Protocol Comparison Matrix

Protocol	Transport	Connection Model	Multiplexed?	Ideal Use Case
HTTP/1.1	TCP	Request-Response (Keep-Alive)	No	Legacy REST APIs
HTTP/2	TCP	Multiplexed binary streams	Yes (Stream-level)	Modern REST APIs, web apps
HTTP/3 (QUIC)	UDP	Multiplexed independent packet streams	Yes (Packet-level)	Mobile devices, spotty cell networks, YouTube
WebSockets	TCP	Persistent, full-duplex bidirectional	No (Single stream)	Real-time chat, live trading, gaming
gRPC	HTTP/2	Typed RPC with binary Protobuf	Yes	Internal microservice-to-microservice traffic

4. Distributed Consensus: Raft & Paxos Essentials

When multiple nodes must agree on a shared state (e.g. Who is the leader? What is the current commit log index?), distributed consensus algorithms ensure agreement even if nodes crash:

Raft in a Nutshell (Leader-Based Consensus)



1. Leader Election:

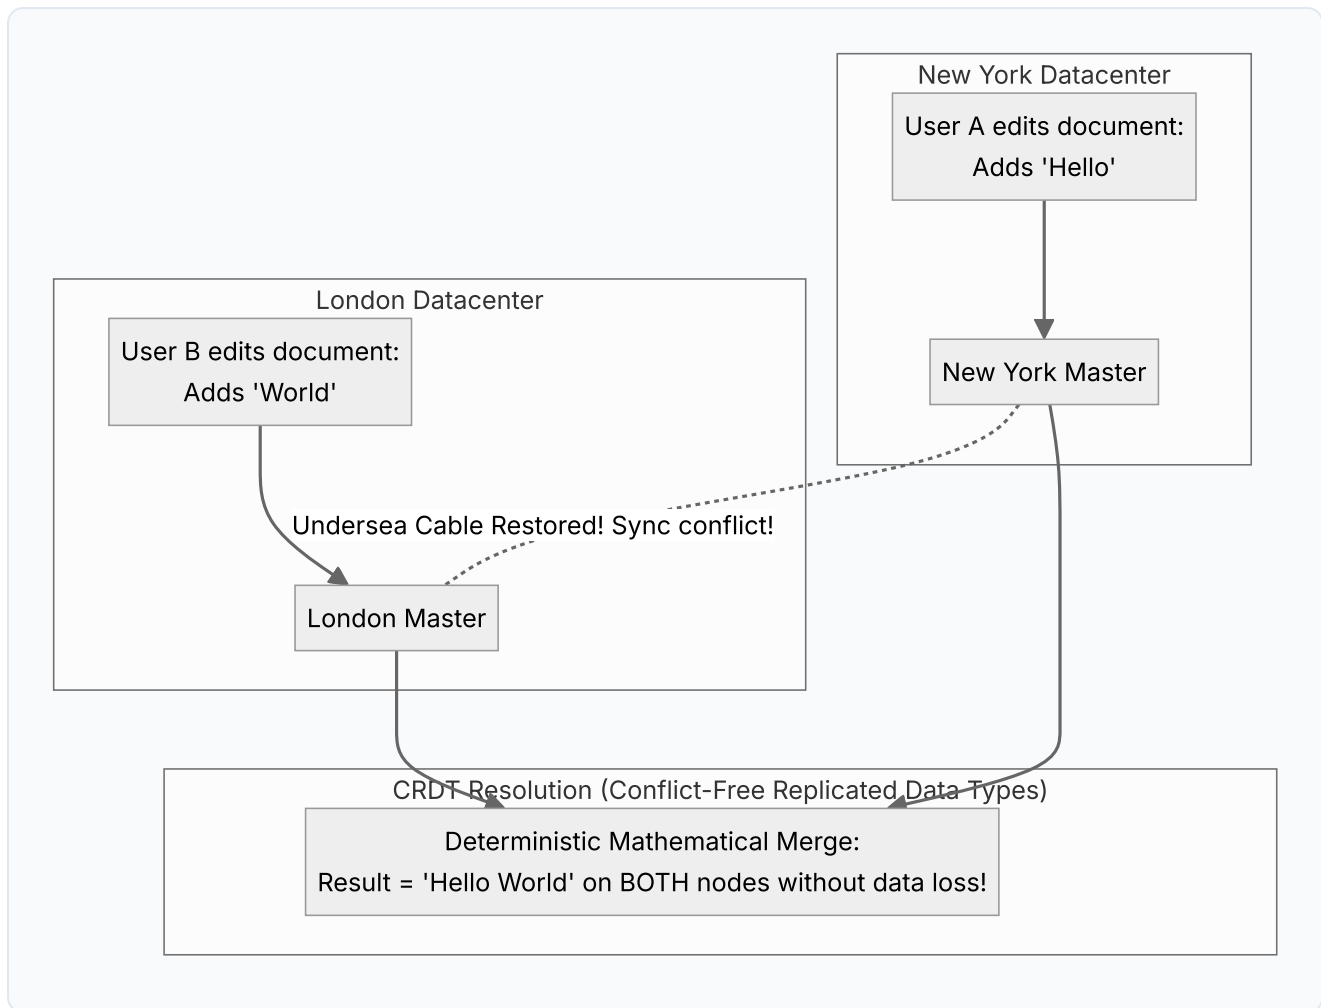
- Nodes start as **Followers**. If a follower hears no heartbeat from a leader within a randomized election timeout (150ms – 300ms), it transitions to **Candidate**, increments the `term` number, and requests votes.
- A candidate receiving votes from a **Quorum** ($\lfloor N/2 \rfloor + 1$) becomes the new **Leader**.
- If two candidates split votes evenly, both timeout with randomized backoffs (150ms vs 280ms), guaranteeing that one triggers an election earlier on the next round!

2. Log Replication:

- Clients send all write requests strictly to the Leader.
- The Leader appends the entry to its local log and broadcasts `AppendEntries` RPCs to all followers.
- Once a Quorum of followers acknowledge the entry, the leader commits it and notifies the client.

5. Active-Active Multi-Region Replication: CRDTs & Vector Clocks

What happens when your database has two primary master databases—one in **New York** and one in **London**—and both accept simultaneous writes while an undersea cable is temporarily cut?



The Conflict Resolution Options:

1. **Last-Write-Wins (LWW):** Compare timestamps and pick the latest write.
 - *The Danger:* Physical computer clocks drift (NTP clock skew). You might accidentally overwrite a critical edit made 5 milliseconds earlier!
2. **Vector Clocks:** Every node maintains an array of logical counters: `[NY: 3, LON: 2]`. This allows the database to detect whether Write A *happened-before* Write B, or if they were **concurrent conflicting writes** requiring application resolution.
3. **CRDTs (Conflict-Free Replicated Data Types):** Data structures mathematically designed so that any concurrent edits can be merged in any arbitrary order, and **all replicas are mathematically guaranteed to converge to the exact same state without locks!** (Used by Figma, Google Docs, Apple Notes, and Redis Enterprise).

Further Reading

- [Raft consensus visualisation](#)
- [Wikipedia: CAP theorem](#)
- [Wikipedia: Vector clock](#)

Check Yourself

Answer in your head or on paper first, then open each answer.

- ▶ 1. State the CAP theorem precisely.
- ▶ 2. What does a quorum read/write with $R+W>N$ guarantee?
- ▶ 3. Why can't you rely on wall clocks to order events across machines?
- ▶ 4. What does Raft do?