


Chapter 09: Load Testing, Benchmarking, & SRE Performance Playbook

The Vanity of Averages *"Our service has an average response time of 45 milliseconds!"* In Site Reliability Engineering (SRE), reporting an average latency is considered malpractice. If 95% of your users experience 10ms response times, but the wealthiest 5% who add 50 items to their cart experience a 12-second timeout, your average might look acceptable while your revenue bleeds dry.

High-performance distributed engineering is defined by **Tail Latency** (*p99* and *p99.9*) and **Saturation Dynamics**. This playbook teaches you how to benchmark distributed services under realistic synthetic loads, diagnose saturation choke points, and design to Google's Four Golden Signals.

1. The Four Golden Signals of Distributed Systems

Google SRE defines four essential metrics for diagnosing distributed systems health:



- 1. Latency:** The time it takes to service a request. Always distinguish between the latency of successful requests vs. failed requests (a failing 500 error that returns in 2ms should not artificially lower your perceived latency!).
- 2. Traffic:** A measure of demand on the system (e.g., HTTP requests/sec, Kafka messages/sec, network I/O bits/sec).
- 3. Errors:** The rate of requests that fail, either explicitly (HTTP 500s) or implicitly (e.g., returning HTTP 200 with an empty JSON payload or incorrect data).

4. **Saturation:** How "full" your service is. Saturation measures constrained resources: CPU quota, RAM, disk I/O, database connection pool occupancy, or socket buffer limits. **Latency degrades exponentially once saturation crosses 80%.**

2. Hands-On Load Testing with Locust


Locust is an industrial-strength, Python-based load-testing framework that allows you to define complex, stateful user journeys in pure code.

The Complete Production Load Test (`locustfile.py`)

Deciphering the Latency Distribution

Notice how `POST /api/v1/checkout` has an **Average of 210ms**, but its *p99* is **2,800ms (2.8 seconds)** and has 120 failures. This tail latency is where customers abandon their carts!

3. The 3 Most Common Saturation Bottlenecks



1. Database Connection Pool Exhaustion

- **Symptom:** Application threads stall waiting on `db.pool.get()`. Latency spikes linearly with traffic. Database CPU itself is low (20%), but application threads are blocked.

- **Root Cause:** Application instances (e.g., 50 pods × 20 pool size = 1,000 connections) exceed Postgres's `max_connections = 300`. Postgres starts rejecting connections.
- **The SRE Fix:** Introduce **PgBouncer** or **AWS RDS Proxy** in transaction pooling mode. This allows 5,000 application threads to multiplex safely over 50 real PostgreSQL server connections.

2. Ephemeral Port Exhaustion (`TIME_WAIT` Buildup)

- **Symptom:** Microservices communicating via HTTP/REST start throwing `Cannot assign requested address` or connection timeouts under 15,000 QPS.
- **Root Cause:** Creating a new HTTP client connection for every request opens a temporary outbound TCP port (range 32768–60999). When closed, TCP sockets enter the `TIME_WAIT` state for 60 seconds (RFC 793) to catch lingering packets. Rapid requests exhaust all ~28,000 available ports.
- **The SRE Fix:** Enforce **HTTP Keep-Alive Connection Pooling** (`urllib3.PoolManager` or `httpx.AsyncClient(limits=Limits(max_keepalive_connections=100))`). Reusing existing TCP sockets eliminates port churn completely.

3. Kubernetes CFS CPU Quota Throttling

- **Symptom:** Pod latency crawls to seconds, yet `kubectl top pod` shows CPU usage is only at 40% of its limit.
- **Root Cause:** The Linux Completely Fair Scheduler (CFS) measures CPU usage in 100ms periods. If a multi-threaded Python/Node app bursts across 8 cores for the first 20ms of the period, it exhausts its allocated quota and the Linux kernel **freezes the entire container for the remaining 80ms!**
- **The SRE Fix:** Remove arbitrary `cpu.limits` on microservices and rely strictly on `cpu.requests` with horizontal pod autoscaling (HPA), or tune the CFS quota period.

4. On-Call Incident Troubleshooting Runbook

When p99 alerts trigger at 2:00 AM, follow this systematic triage loop:

1. **Check Error Codes:** Are failures 502/504 (Upstream timeout / Bad Gateway) or 500 (Application crash)?
2. **Check Database Metrics:**
 - Active connections vs. Max connections pool.
 - Longest running active queries (`SELECT pid, now() - query_start, query FROM pg_stat_activity WHERE state ≠ 'idle'`).
3. **Check Container Health:**

- Is Pod restarting due to `OOMKilled` (Exit code 137)?
- Is container CPU throttled (`container_cpu_cfs_throttled_periods_total`)?

4. **Isolate Recent Changes:** Did a deployment or database schema migration occur in the last 60 minutes? If yes, **rollback first, ask questions later.**

Further Reading

- [Google SRE book: service level objectives](#)
 - [Gil Tene: How NOT to measure latency](#)
 - [k6 documentation](#)
-

Check Yourself

Answer in your head or on paper first, then open each answer.

- ▶ 1. Why measure percentiles instead of averages?
- ▶ 2. What is coordinated omission?
- ▶ 3. What is an SLO versus an SLA?
- ▶ 4. What is an error budget?