

HLD Chapter 11: Scenario-Based HLD Crisis Scenarios & Senior Staff Grills

Core Learning Objective: Master real-world distributed systems outage scenarios and high-pressure interviewer follow-ups. Learn how to diagnose Cassandra zombie data resurrections, Redis 100% CPU lockouts, NTP clock drift in Snowflake IDs, and multi-region failover.

1. Production Crisis Scenarios

Scenario 1: The Cassandra "Zombie Data" Resurrect Bug

The Interviewer Asks:

"In our chat platform built on Apache Cassandra, a user deletes a message. Three weeks later, the deleted message mysteriously reappears in their chat history! What happened under the hood, and how do you prevent it?"

The Staff-Level Answer:

- **Root Cause (Tombstone Eviction & Compaction Lag):**
 - In Cassandra (LSM-Tree), a `DELETE` does not immediately erase data from disk; it writes a marker called a **Tombstone** with a deletion timestamp.
 - Cassandra keeps tombstones for a configurable window (`gc_grace_seconds` , default 10 days) to allow all replicas to receive the deletion.
 - If Replica Node 3 was offline for 12 days (exceeding `gc_grace_seconds`), Replicas 1 and 2 compacted and evicted the tombstone.
 - When Node 3 rejoined the cluster, its stale SSTables still held the original undeleted message! When a read request queries Node 3, Cassandra treats the older message as live data and "resurrects" it into the other replicas.
- **The Fix:**
 - Run **Incremental / Scheduled Repairs** (`nodetool repair`) frequently within the `gc_grace_seconds` window.
 - Never allow an offline node to rejoin the cluster if it has been down longer than `gc_grace_seconds` ; decommission and re-bootstrap it with a clean snapshot.

Scenario 2: Redis Cluster at 100% CPU Utilization

The Interviewer Asks:

"Our primary Redis cluster CPU suddenly spikes to 100%, causing request timeouts across the entire platform. Redis is single-threaded for command execution. How do you diagnose and mitigate this live in production?"

The Staff-Level Answer:

1. Diagnosis:

- Run `redis-cli --bigkeys` and `redis-cli --hotkeys` to detect if a single key is consuming massive bandwidth or receiving 50,000 QPS.
- Run `SLOWLOG GET 10` to check for $O(N)$ operations accidentally executed in production (e.g. `KEYS *` instead of `SCAN`, or massive `HGETALL` / `SMEMBERS` on sets with 100,000 elements).

2. Mitigations:

- **Hot Key Partitioning (Key Sharding):** If `featured_products` is hammered by all users, replicate the key across K sub-keys: `featured_products_1`, `featured_products_2`, ..., `featured_products_N`. Application servers pick a random sub-key, spreading read load evenly across Redis shards.
- **Near-Cache (L1 In-Memory Cache):** Add a tiny in-process memory cache (e.g. `cachetools` or Go `sync.Map` with 10-second TTL) directly on the API Gateway instances, reducing Redis queries by 90%.

Scenario 3: Clock Drift & Snowflake ID Collisions

The Interviewer Asks:

"In our distributed Snowflake ID generator, a server's local clock synchronizes via NTP and jumps backward by 250 milliseconds. What happens to ID generation, and how do we prevent duplicate primary keys?"

The Staff-Level Answer:

• The Problem:


Snowflake embeds a 41-bit millisecond timestamp into the ID. If the system clock jumps

backward, the node risks generating duplicate IDs that were already issued during the prior physical milliseconds.

- **The Production Solution:**

- Track the `last_timestamp` generated on each node.
- If `current_timestamp < last_timestamp`:
 - If the drift is small ($< 10\text{ms}$): Sleep the thread until the clock catches up to `last_timestamp`.
 - If the drift is significant ($> 10\text{ms}$): Refuse to issue IDs, increment a clock drift failure metric, and failover ID generation to another healthy node.

2. Multi-Region Active-Active Replication



Handling Multi-Region Conflicts (Active-Active)

If a user updates their profile in US-East while their automated bot updates it in EU-West simultaneously:

1. **Last-Write-Wins (LWW):** Uses NTP timestamp to pick the highest clock value (Danger: Clock skew can overwrite newer changes).
2. **Conflict-Free Replicated Data Types (CRDTs):** Mathematically convergent data structures that merge concurrent updates without centralized coordination.

Further Reading

- [System Design Primer](#)
- [Stripe: rate limiters](#)

- [AWS Builders' Library](#)
-

Check Yourself

Answer in your head or on paper first, then open each answer.

- ▶ 1. How do you start 'design a rate limiter'?
- ▶ 2. Strong versus eventual consistency: how do you justify choosing?
- ▶ 3. How do you handle 'it must be highly available'?
- ▶ 4. What closes a strong answer?