

HLD Case Study 5: Planetary Video Streaming (YouTube / Netflix)

Key Focus Areas: Upload and transcoding pipelines, chunking DAGs, Adaptive Bitrate Streaming (HLS/DASH), CDN and origin shielding, storage and egress cost.

1. Requirements and Scope

Functional: upload large videos, process into multiple resolutions and codecs, stream with adaptive quality, search/metadata, view counts, resumable playback. **Non-functional:** start playback in under 2 seconds p95; no rebuffering on normal networks; uploads are durable and resumable; processed video available within minutes for typical files; very high read-to-write ratio. **Out of scope:** recommendations, live streaming (mention differences), DRM details, ads.

2. Estimates

Assume 2B MAU, 800M daily viewers, average 40 minutes watched/day, 5 Mbps average delivered bitrate; 500 hours of video uploaded per minute (YouTube-like).

Takeaways: delivery is **hundreds of Tbps at peak**, so the architecture is a CDN problem, not an application-server problem; ingest is about 1.4 PB/day of raw video, and transcoding CPU/GPU cost dominates the write path.


3. API Design

- `POST /uploads` returns a resumable upload session (chunked, e.g. 8 MB parts); `PUT` parts to object storage via presigned URLs; `POST /uploads/{id}/complete`.
- `GET /videos/{id}` returns metadata and the **manifest URL** (`.m3u8` for HLS or `.mpd` for DASH).
- `GET /videos/{id}/segments/ ...` is served by the CDN, not the API.
- `POST /videos/{id}/view` (batched, eventually consistent counters).

4. Data Model

- **videos** (metadata DB, sharded by `video_id`): owner, title, status `UPLOADING|PROCESSING|READY|FAILED`, duration, renditions list.
- **object storage layout:** `raw/{video_id}/source`, `proc/{video_id}/{rendition}/seg_00001.m4s`, `manifest.m3u8`.
- **view counts:** counters in Redis/streaming aggregation (Kafka to Flink) flushed to a store; counts are approximate in real time.

5. Architecture



6. Deep Dive: Transcoding and Adaptive Bitrate

Chunked parallel transcoding. A 2-hour movie encoded sequentially takes hours. Instead, split at keyframe (GOP) boundaries into segments, encode all segments of all renditions in parallel across

a fleet, then assemble manifests. A DAG orchestrator tracks per-segment tasks with retries; a failed segment is re-encoded alone.

Bitrate ladder. Example rungs: 240p at 0.4 Mbps, 360p at 0.8, 480p at 1.5, 720p at 3, 1080p at 5.5, 4K at 16. Modern services use **per-title encoding**: simple animation needs fewer bits than a sports match at the same resolution, so the ladder is tuned per video to save bandwidth at equal quality. Codecs: H.264 (universal), VP9/AV1 (30 to 50% fewer bits, costlier to encode, encode only popular videos in AV1).

Adaptive bitrate (ABR). The player downloads the manifest, picks a rendition based on measured throughput and buffer level, and re-evaluates before **each segment**. If throughput drops it steps down; when the buffer is healthy it steps up. Segment length is a trade-off: short segments (2 s) adapt faster but add request overhead; long (6 s) are efficient but slow to react.

CDN and origin shield. Edge caches serve the hot head; misses go to a regional **origin shield** which collapses duplicate misses before they reach object storage. Because video is immutable and segments are static files, cache hit ratios above 95% are achievable. Long-tail videos (rarely watched) are served from shield/origin and **not pre-encoded in every codec** (encode lazily on demand to save storage).

7. Scaling and Bottlenecks

- **Egress bandwidth:** peering and ISP-embedded caches (Netflix Open Connect style) move bytes close to users and cut transit cost.
- **Transcoding cost:** priority queues (new popular channels first), spot/preemptible compute for non-urgent work, hardware encoders for high-volume codecs.
- **Viral video:** pre-warm CDN for videos with fast-growing views; request coalescing at the shield.
- **Storage:** hot/warm/cold tiers; delete unused renditions after a period; deduplicate re-uploads by content hash.
- **View counting:** do not write per view; buffer in stream processing, dedupe by (user, video, window).

8. Failure Modes and Mitigations

Failure	Mitigation
Upload interrupted	resumable parts with checksums; client resumes from last part
Transcoder crash	per-segment retries; DAG state persisted; idempotent output paths
Corrupt source	validation step, mark FAILED with reason
CDN POP outage	DNS/Anycast steers to next POP; player retries another CDN
Origin overload on a cold cache	origin shield + request coalescing + rate limits
Partial manifest (some renditions missing)	publish manifest only after the minimum ladder is complete; add higher rungs later

9. Trade-offs

- **Pre-transcode all vs on-demand:** pre-transcoding costs storage and compute for videos nobody watches; on-demand adds first-view latency. Mixed: pre-encode a base ladder, add expensive codecs after popularity crosses a threshold.
- **HLS vs DASH:** HLS has universal Apple support; DASH is codec-flexible. CMAF fragments let both share the same media files.
- **Segment size:** latency vs overhead, as above.
- **Client-side ABR vs server-side:** client ABR scales for free; server control gives better fairness but adds state.

10. Interview Timeline and Follow-ups

Spend most time on the upload to transcode to CDN pipeline and ABR; mention cost levers last. **Follow-ups:** How would live streaming differ? (low-latency ingest via RTMP/WebRTC, 1 to 2 s segments or chunked CMAF, no pre-encoding, shorter caches). How do you resume playback across devices? (store last position per user in a KV store, updated every N seconds). How do you handle copyrighted content? (fingerprinting at ingest). How do you reduce start-up delay? (fetch a low first rung, prefetch next segments, connection reuse/HTTP/3).