

System 11: Distributed Cache (Redis Cluster Architecture)

Zero-Prerequisite Intuition: The "Sticky Note on the Monitor" Metaphor What is a cache, and why do we distribute it? Imagine you work as an accountant. Every time a client asks for their tax ID, you have to stand up, walk 50 meters down a long hallway, unlock a giant iron fireproof vault, search through 10,000 paper folders, find the document, walk 50 meters back, and tell the client. That long walk to the vault takes **5 minutes**. That is your **Database**.

If 500 clients call you in one hour, you will spend your entire day running up and down the hallway until you collapse from exhaustion.

What do you do instead? The first time a client asks for their tax ID, you write it on a small **Sticky Note** and slap it right onto the bezel of your computer monitor. The next 499 times that client asks, you don't walk down the hallway—you glance at the sticky note in **0.1 seconds**. That is a **Cache**.

But what happens when you have 10 million clients? You cannot fit 10 million sticky notes on one monitor! You need an entire room of accountants, each holding a fraction of the sticky notes. That is a **Distributed Cache**.

1. System Scale & Capacity Estimation

Let's design a distributed cache capable of supporting Twitter/Meta-scale read traffic:

- **Total In-Memory Data:** 10 Terabytes of active keys.
- **Read Throughput:** 1,000,000 Queries Per Second (QPS).
- **Write Throughput:** 100,000 QPS.
- **Target Latency:** $p99 < 1.0$ millisecond.

Hardware Cluster Sizing:

- High-memory cloud instances (e.g. AWS `r6i.4xlarge`): 128 GB RAM per instance.
- To safely store 10 TB with a 30% headroom buffer for key fragmentation and replication:

$$\text{RAM Needed} = 10 \text{ TB} \times 1.30 = 13 \text{ TB}$$

$$\text{Nodes Required} = \frac{13,000 \text{ GB}}{128 \text{ GB/node}} \approx 102 \text{ Master Nodes}$$

- With 1 Replica per Master for High Availability: **204 total nodes**.


2. Distributed Architecture: Hash Slots & Gossip Protocol

How do 102 independent Redis nodes decide which server holds which sticky note without a centralized master bottleneck?

The 16,384 Hash Slot Algorithm

Redis Cluster does not use simple modulo ($\text{hash}(\text{key}) \% N$) because adding or removing a node would invalidate 99% of all cached keys! Instead, it divides the universe into **16,384 virtual Hash Slots**:

$$\text{Slot} = \text{CRC16}(\text{key}) \pmod{16384}$$



1. **Node A** is assigned slots 0 to 5460 .
2. **Node B** is assigned slots 5461 to 10922 .
3. **Node C** is assigned slots 10923 to 16383 .

4. When a new node joins, slots are migrated incrementally without taking the cluster offline!

Smart Clients & The **MOVED** Redirect

Redis clients cache a local routing table of which node owns which slot. If a client mistakenly sends a query for Slot 8100 to Node A:

- Node A does not proxy the query (proxying adds latency).
- Node A returns an instant error: `-MOVED 8100 10.0.0.2:6379`.
- The client updates its local slot cache and queries Node B directly!


3. Node Failure & Consensus via Gossip Protocol

Every node connects to every other node via a secondary **Cluster Bus port** (standard port + 10,000, e.g., 16379).

- Nodes exchange **PING / PONG heartbeat packets** using a decentralized **Gossip Protocol**.
- If Node A stops responding to Node B for `cluster-node-timeout` (e.g. 5 seconds), Node B flags Node A as **PFAIL** (Possible Failure).
- Once a majority of master nodes agree that Node A is unreachable, **PFAIL** transitions to **FAIL**.
- The replica node of Node A initiates an automated election, promotes itself to Master, and claims the hash slots without downtime!

4. Cache Invalidation & Eviction Algorithms

When the 128 GB RAM on a node fills up completely, what happens when new data arrives?



Approximated LRU vs True LRU

A true Least-Recently-Used (LRU) cache requires maintaining a doubly-linked list of every key in memory. In a 10 TB cluster, the pointers alone would waste 2 TB of RAM!

- **Redis Solution:** Redis samples 5 random keys at random and evicts the one with the oldest access time (`lru_clock`). This statistical approximation achieves 99% of true LRU efficiency with **zero memory overhead!**

5. Summary Architecture Blueprint

Component	Technology / Strategy	Production Decision Justification
Partitioning Scheme	16,384 Hash Slots with CRC16	Minimizes re-sharding data movement when adding/removing cluster nodes
Failover Mechanism	Raft-like Quorum Voting over Gossip	Eliminates single-point-of-failure centralized coordinators (like ZooKeeper)
Replication	Asynchronous Master-to-Replica	Guarantees sub-millisecond write latency ($p99 < 1ms$) by not waiting for cross-network replica acks
Eviction Policy	<code>allkeys-lru</code> (Approximated 5-sample)	Protects RAM from unbounded growth while prioritizing popular hot items