

HLD Case Study 18: Typeahead / Search Autocomplete

Key Focus Areas: Prefix lookup in memory, top-k at every node, offline versus online ranking, sharding by prefix, latency budget, trending updates.

1. Requirements and Scope

Functional: as the user types, return the 5 best completions for the current prefix, ranked by popularity (and recency); suggestions follow new trends within minutes to hours; offensive or blocked terms never appear. **Non-functional:** end-to-end under 100 ms (the user is typing); the service is read-heavy and tolerates **slightly stale** rankings; high availability, because a missing suggestion box degrades but does not break search. **Out of scope for this answer:** spelling correction, personalised ranking beyond a short overlay, multi-language segmentation.

2. Back-of-the-Envelope Estimates

The whole index fits in the RAM of one large machine, so the problem is **serving 300,000 to 900,000 lookups per second with low latency**, not storage. Replicate the index many times rather than sharding it deeply.

3. API Design


Cache headers matter: suggestions for popular prefixes can be cached for 1 to 5 minutes at the CDN and in the browser. Debounce on the client (wait about 100 to 200 ms after the last keystroke) and cancel superseded requests.

4. Data Model

Offline pipeline (batch or stream): `search logs → count by (term, time window) → filter → build index`. Serving index: a trie in memory where **each node stores the top-k completions for its prefix**, so a lookup is "walk the prefix, return the node's list" with no subtree scan.

Item	Stored where	Notes
Raw query logs	Kafka then object storage	The source of truth for rankings
Term counts per window	Batch job output	Decay old counts so trends can rise
Serving trie (immutable snapshot)	RAM on every serving node	Rebuilt and swapped atomically
Block list	Small side set	Applied at build time and again at serve time

5. Architecture



6. Deep Dive: The Trie with Top-k at Each Node

Walking a trie and collecting every completion under a prefix is too slow for short prefixes ("a" has millions). Precompute the best k at every node:

Cost. Memory grows with `terms x prefix length x k`, which is why real systems share prefixes (a compressed or radix trie), cap k at 5 to 10, and drop rare terms. Lookups cost `O(prefix length)` regardless of how many terms exist.

Ranking signals. Start with frequency over a time window; add recency decay (a weekly count weighted above a yearly count), geography and language, and a click-through signal ("which suggestion did users actually pick"). Keep ranking offline; the serving node only reads lists.

Trending. An hourly rebuild is too slow for breaking news. Run a small streaming job that detects terms whose rate jumped, and merge those into the results as an overlay at serve time.

7. Scaling and Bottlenecks

1. **Replicate, do not shard, the small index:** a few gigabytes per node means every node can answer every prefix; scale by adding replicas behind a load balancer.
 2. **Shard only if the index outgrows memory:** shard by the first one or two characters, and remember skew (the letter "s" is far hotter than "x"); split hot prefixes further.
 3. **Cache aggressively:** a few thousand short prefixes account for most traffic; CDN and in-process caches absorb them.
 4. **Snapshot swap:** build the new trie off to the side, then switch a pointer atomically so lookups never see a half-built index.
 5. **Client discipline:** debounce, cancel in-flight requests, and do not query below a minimum length for very common prefixes if latency demands it.
-

8. Failure Modes and Mitigations

Failure	Effect	Mitigation
Suggest service down	No suggestions, search still works	The UI hides the dropdown; never block the search box
Bad build (empty or corrupt trie)	Wrong or empty suggestions everywhere	Validate a snapshot (size, sample queries) before swapping; keep the previous snapshot for instant rollback
Offensive term trending	Reputation damage	Block list applied at build and serve time, plus a rapid takedown path
Stale trends	Users see yesterday's popular terms	Streaming overlay and short snapshot intervals
Hot prefix overload	One node saturates	Replicas plus CDN caching for the top prefixes

9. Trade-offs and Alternatives

- **Trie versus a search index (Elasticsearch completion suggester):** a search engine is faster to build and richer (fuzzy matching, filters); a custom in-memory trie wins on raw latency and cost at very high QPS.
 - **Per-node top-k versus computing on the fly:** per-node lists cost memory and make lookups constant time; on-the-fly scans cost nothing in memory and are too slow for short prefixes.
 - **Freshness versus stability:** rebuilding often tracks trends but makes suggestions jitter; users prefer stable lists with a small trending slice.
 - **Personalisation:** a small per-user overlay (recent searches) merged client-side or at the edge keeps the shared index cacheable.
-

10. Interview Timeline (45 minutes) and Follow-ups

Minutes	Do
0 to 5	Latency target, freshness expectation, scope
5 to 10	Estimates: QPS and index size (it fits in memory)
10 to 20	Data structure: trie with top-k per node, and why
20 to 30	Offline pipeline, serving architecture, snapshot swap
30 to 40	Trending overlay, ranking, sharding and caching
40 to 45	Failure modes, abuse filtering, personalisation

Follow-ups to prepare: How do you update the trie without downtime? How do you handle a term that suddenly trends? How do you shard when the index no longer fits one node? How do you avoid suggesting private or offensive queries? How do you measure suggestion quality?