

Chapter 13: Vector Database Internals & HNSW Graph Math

Zero-Prerequisite Intuition: The "Galaxy Star Map" Metaphor What is a vector embedding, and why do we need special Vector Databases (Pinecone, Qdrant, Milvus, Chroma)? Imagine looking up at the night sky. Every star has 3D coordinates in space: (X, Y, Z) .

- Earth is at $(0, 0, 0)$.
- Mars is nearby at $(0.5, 0.2, 0.1)$.
- The Andromeda Galaxy is far away at $(2500000, 890000, 120000)$. Because Earth and Mars have coordinates that are close to each other, you know they are in the same neighborhood!

A **Vector Embedding Model** does the exact same thing for human thoughts and text. It reads a sentence and assigns it coordinates in a **1,536-dimensional space**:

- "The dog barked" $\rightarrow [0.24, -0.81, 0.12, \dots]$
- "A puppy made a sound" $\rightarrow [0.26, -0.79, 0.11, \dots]$ (Extremely close coordinates!)
- "Stock market interest rates rose" $\rightarrow [-0.91, 0.34, -0.88, \dots]$ (Far away coordinates!)

But here is the crisis: If you have **10,000,000 document vectors**, computing the exact distance between your search query and all 10 million vectors requires billions of floating-point operations. Your search takes **6 seconds**!

A **Vector Database** uses **Approximate Nearest Neighbors (ANN)** to search 10 million vectors in **less than 8 milliseconds**!

1. Distance Metrics: Cosine vs. Dot Product vs. Euclidean (\$L2\$)

How do we mathematically measure the "distance" between two thought vectors $\mathbf{u}$ and $\mathbf{v}$?

```
graph LR
    subgraph Distance_Calculations ["Mathematical Similarity Metrics"]
        L2["Euclidean (L2): Direct Straight-Line Distance<br>d(u, v) = √(Σ(u_i - v_i)²)"]
        Dot["Dot Product: Magnitude × Angle<br>u · v = Σ(u_i × v_i)"]
        Cos["Cosine Similarity: Pure Angular Direction<br>cos(θ) = (u · v) / (||u|| × ||v||)"]
    end
```

- **Euclidean (\$L2\$) Distance:** Measures the geometric straight-line physical distance. Sensitive to vector magnitude (document length).
- **Dot Product:** Fast to compute on modern CPU/GPU hardware using SIMD instructions.
- **Cosine Similarity:** Measures the angle θ between vectors, completely ignoring document length. $\cos(\theta) = \frac{\mathbf{u} \cdot \mathbf{v}}{|\mathbf{u}| |\mathbf{v}|}$

TIP

Production Optimization: If you normalize all embedding vectors to unit length ($|\mathbf{u}| = 1$) upon ingestion, **Cosine Similarity becomes mathematically identical to Dot Product!** You can replace expensive divisions and square roots with a single blazing-fast hardware vector dot product.

2. The HNSW Algorithm: Multi-Layer Skip-Graphs

The undisputed state-of-the-art algorithm for vector search is **Hierarchical Navigable Small World (HNSW)**.

The "Express Train vs. Local Walking" Metaphor

How do you travel from New York to a specific house in Los Angeles?

- You don't walk street-by-street across 3,000 miles (that is brute force $O(N)$).
- You take an **Express Airplane** across the continent (Layer 2).
- You land at LAX and take a **Regional Highway** to the neighborhood (Layer 1).
- You take a **Local Street** to the exact doorstep (Layer 0).

```
flowchart TD
```

```
    subgraph Layer_2 ["Layer 2: Top Expressway (Sparse Graph, Long Links)"]
```

```
        L2_A["Node 10"] —→ L2_B["Node 800"]
```

```
    end
```

```
    subgraph Layer_1 ["Layer 1: Regional Roads (Medium Density)"]
```

```
        L1_A["Node 10"] → L1_C["Node 250"] → L1_B["Node 800"]
```

```
    end
```

```
    subgraph Layer_0 ["Layer 0: Local Streets (All Vectors Present)"]
```

```
        L0_A["Node 10"] → L0_D["Node 45"] → L0_C["Node 250"] → L0_E["Node 720"] →
```

```
    end
```

```
    Query["Search Query Vector Q"] -. "1. Start at Top Layer" ·-> L2_A
```

```
    L2_A -. "2. Greedy Hop to Closest Node" ·-> L2_B
```

```
    L2_B -. "3. Drop Down" ·-> L1_B
```

```
    L1_B -. "4. Drop Down to Layer 0" ·-> L0_B
```

```
    L0_B -. "5. Explore Local Neighbors" ·-> L0_E
```


The Search Mechanics:

1. Start at the top layer entry point with greedy routing: jump to whichever neighboring node is closest to query vector Q .
 2. When no neighbor is closer to Q , drop down one layer and resume greedy routing from that node.
 3. At Layer 0 (the bottom layer containing all vectors), execute a **Beam Search** with size `efSearch` to extract the Top- K nearest neighbors.
- **Complexity:** Scales in logarithmic time: $\mathbf{O}(\log N)$ rather than linear time $\mathbf{O}(N)$! Searching 10 million vectors takes only ~20 graph hops!

3. Product Quantization (IVF-PQ): Compressing Vectors

Storing 10 million vectors of 1,536 floats in RAM requires: $10,000,000 \times 1536 \times 4 \text{ bytes} \approx 61.4 \text{ Gigabytes of RAM!}$ RAM is expensive (\$100s/month).

Product Quantization (PQ) compresses this by 95% down to **3 GB** using lossy vector quantization:



- 1. Divide the 1,536-dimensional vector into 64 sub-vectors of 24 dimensions each.
- 2. Run k -means clustering on each sub-space to generate 256 centroid prototypes.
- 3. Instead of storing 24 floats (96 bytes), store a single **8-bit integer (1 byte)** representing the centroid index!
- 4. Querying computes distance using pre-computed centroid lookup tables, speeding up distance computations by $10\times$.

4. Metadata Filtering: Pre-Filtering vs. Single-Stage Hybrid Search

In real applications, you rarely search vectors in isolation. You search: *"Find documents similar to 'billing refund' WHERE `user_id = 42` AND `created_year ≥ 2025`"*.

Filtering Strategy	How It Works	Danger / Failure Mode
Post-Filtering	Run HNSW vector search to find Top-100 nearest items, then discard items that don't match <code>user_id = 42</code> .	Zero Results Disaster: If only 1 out of 10,000 items belongs to user 42, the top-100 will contain 0 matching items, returning an empty list!
Pre-Filtering	Query SQL index first to find all IDs matching <code>user_id = 42</code> , then run brute force vector distance only on those matching IDs.	Inefficient if user 42 has 500,000 items (falls back to slow brute force).
Single-Stage Filtered HNSW (Qdrant / Milvus)	Navigates the HNSW graph while skipping graph nodes during beam search that fail the metadata bitset.	Optimal: Guarantees exact Top- K while preserving $O(\log N)$ graph search speeds.