

Chapter 14: Model Fine-Tuning, LoRA, & DPO Alignment


Zero-Prerequisite Intuition: The "College Graduate vs. Specialized Surgeon" Metaphor

What is Fine-Tuning, and how is it different from RAG and Prompt Engineering? Imagine hiring a brilliant college graduate who read the entire world's encyclopedias and libraries.

- **Prompt Engineering:** You give the graduate an instruction: *"Please read this email and summarize it in three bullet points."*
- **RAG (Retrieval-Augmented Generation):** The graduate doesn't know your company's private financial records. So you hand them a physical binder containing yesterday's ledger: *"Answer the client's question using ONLY page 4 of this binder."*
- **Fine-Tuning:** You want this graduate to become a **Cardiovascular Surgeon**. Handing them a medical textbook during an open-heart surgery is too slow—they must absorb the instincts, surgical vocabulary, muscle memory, and precise medical formatting into their subconscious mind!

Fine-tuning alters the internal weights (the brain synapses) of the model. This chapter covers the definitive decision framework for when to fine-tune vs. RAG, the mathematics of **LoRA** and **QLoRA**, and modern model alignment with **DPO (Direct Preference Optimization)**.

1. The Decision Matrix: Prompting vs. RAG vs. Fine-Tuning



Dimension	Prompt Engineering	RAG	Fine-Tuning (LoRA)
Primary Purpose	Guiding instruction behavior	Injecting dynamic external facts	Modifying style, syntax, and instincts
New Knowledge	Zero (limited by context window)	High (retrieves from billions of docs)	Low (prone to catastrophic forgetting)
Hallucination Risk	Moderate	Very Low (Grounded in context)	High (if used to memorize facts)
Latency & Cost	Low setup; High token costs	Medium latency (vector search hop)	Zero retrieval latency; Low token count

2. Parameter-Efficient Fine-Tuning: The LoRA Mathematics

Why can't we just train all weights of a modern Large Language Model (e.g., Llama-3 70B)? To train a 70B parameter model in standard 16-bit precision:

- Weights: 140 GB.
- Gradients: 140 GB.
- Adam Optimizer States (Momentum & Variance): **560 GB**.
- Total VRAM required: **Over 840 Gigabytes of GPU VRAM** (Requiring a cluster of $12 \times \$35,000$ NVIDIA H100 GPUs)!

Low-Rank Adaptation (LoRA)

Published by Microsoft researchers (Hu et al.), LoRA realizes that weight updates during adaptation have a very low **intrinsic rank**.


Instead of updating the full $d \times k$ weight matrix W_0 :

$$W = W_0 + \Delta W$$

LoRA **freezes the pre-trained weights W_0 completely** and decomposes the update ΔW into two tiny low-rank matrices A and B :

$$\Delta W = B \times A$$

Where $W_0 \in \mathbb{R}^{d \times k}$, $B \in \mathbb{R}^{d \times r}$, and $A \in \mathbb{R}^{r \times k}$, with rank $r \ll \min(d, k)$ (typically $r = 8$ or 16).



- **Parameters Reduced by 99.9%:** For a 4096×4096 matrix, full training updates **16,777,216 parameters**. With LoRA rank $r = 8$, matrices A and B have only $4096 \times 8 + 8 \times 4096 = 65,536$ parameters!
- **Zero Inference Latency Overhead:** At deployment, you can mathematically merge the adapter weights directly into the base model: $W_{\text{final}} = W_0 + \frac{\alpha}{r}(BA)$.

3. QLoRA: Fine-Tuning on a Single Consumer GPU


Tim Dettmers et al. introduced **QLoRA (Quantized LoRA)**, enabling 70B parameter models to be fine-tuned on a single 48 GB GPU.

The 3 Core Innovations of QLoRA:

1. **NF4 (NormalFloat 4-bit) Quantization:** A mathematically optimal information-theoretic data type for normally distributed neural network weights.
2. **Double Quantization (DQ):** Quantizes the quantization constants themselves, saving an additional 0.37 bits per parameter.
3. **Paged Optimizers:** Uses CUDA Unified Memory to automatically page optimizer memory spikes between GPU VRAM and CPU RAM, preventing out-of-memory crashes.

4. Alignment & Preference Optimization: DPO vs. RLHF

Once a model is fine-tuned, how do we prevent it from being toxic, and train it to prefer helpful, concise answers?



The DPO Breakthrough (Rafailov et al., Stanford)

Traditional Reinforcement Learning from Human Feedback (RLHF) required training a separate Reward Model, then using unstable PPO reinforcement learning loops that constantly diverged or suffered mode collapse.

DPO mathematically proves that the reward function can be expressed directly through the language model's own probabilities:

$$\mathcal{L}_{\text{DPO}}(\pi_\theta; \pi_{\text{ref}}) = -\mathbb{E}_{(x, y_w, y_l)} \left[\log \sigma \left(\beta \log \frac{\pi_\theta(y_w|x)}{\pi_{\text{ref}}(y_w|x)} - \beta \log \frac{\pi_\theta(y_l|x)}{\pi_{\text{ref}}(y_l|x)} \right) \right]$$

Where:

- x is the prompt.
- y_w is the **winning / preferred** human response.
- y_l is the **losing / rejected** response.
- π_θ is the active model, and π_{ref} is the frozen baseline model.

DPO increases the probability of preferred answers while decreasing the probability of rejected answers in a **single standard cross-entropy training pass!**

Further Reading

- [LoRA paper](#)
- [DPO paper](#)
- [Hugging Face PEFT](#)

Check Yourself

Answer in your head or on paper first, then open each answer.

- ▶ 1. What does LoRA train?
- ▶ 2. What does DPO optimise?
- ▶ 3. When fine-tune instead of using RAG?
- ▶ 4. Why hold out an evaluation set before fine-tuning?